

DocName:pipeline_walkthrough.md

What we learned from this pipeline walkthrough

1. Why naïve RAG breaks on regulatory docs

- **Structure > text:** Tax docs encode meaning via hierarchy, cross-references, and scope—not just sentences.
 - **Lexical similarity is misleading:** “Box 2e” vs “Box 2f” look close but mean different things.
 - **Negation flips truth:** Without polarity, embeddings create false positives.
👉 **Conclusion:** You must model structure and semantics explicitly.
-

2. Knowledge graphs are the right abstraction

- Anchors (boxes/sections) act as **ownership boundaries**.
 - Typed edges encode **how** concepts relate (includes, excludes, applies_if), not just that they co-occur.
 - Polarity (positive vs negative) is essential to prevent incorrect associations.
👉 **Conclusion:** A graph captures what embeddings alone cannot.
-

3. Physical layout understanding is non-negotiable

- Font size, bolding, columns, and page flow determine meaning.
 - Column-aware ordering prevents misassignment across columns.
 - Body-font inference anchors all downstream classification.
👉 **Conclusion:** Semantics depend on accurate physical reconstruction.
-

4. Regex-first + heuristics beats LLM-first here

- Deterministic regex + position rules give **high precision** for roles and edges.
 - Explicit uncertainty (`role=NULL`) is better than hallucinated labels.
 - LLMs are better as **judges**, not primary extractors, at this stage.
👉 **Conclusion:** Use LLMs to validate, not to guess structure.
-

5. Sentence-level gating is critical

- Extracting edges per sentence avoids cross-sentence bleed.
 - Subset cues (“any part of”, “portion of”) must be explicit.
 - Narration filters prevent descriptive text from becoming rules.
- 👉 Conclusion: Precision requires aggressive gating.
-

6. Negation must dominate precedence

- **excludes** edges override all others.
 - Negative knowledge creates **hard negatives** for training.
 - Without this, models learn the wrong associations.
- 👉 Conclusion: Negation is first-class knowledge.
-

7. Validation is as important as extraction

- Deterministic checks catch structural and statistical failures early.
 - Edge distribution (A8) revealed a real weakness despite “working” output.
 - Provenance and DAG integrity are mandatory for trust and traversal.
- 👉 Conclusion: If it isn’t validated, it isn’t usable.
-

8. Training data quality comes from graphs, not chunks

- Graph edges naturally generate **balanced, labeled contrastive pairs**.
 - **excludes** edges enable high-quality hard negatives.
 - Stratified sampling prevents reference spam from dominating.
- 👉 Conclusion: Better graphs → better training → better RAG.
-

9. Current bottleneck is semantic recall, not structure

- Structure and anchors are solid (22/22 boxes, clean DAG).
 - Failure is due to **under-extraction of semantic edges**, not noise.
 - Improving typed-edge patterns will likely flip the system to PASS.
- 👉 Conclusion: The system is directionally correct; extraction depth is the next lever.

10. Big picture takeaway

This pipeline shows that **high-stakes documents require symbolic structure + semantic polarity**, not just vector similarity.

You're effectively building a **document compiler**—turning PDFs into executable knowledge.

If you want, next we can:

- Design new patterns to boost semantic edge recall
- Add lightweight ML to assist typed-edge detection
- Define retrieval strategies that exploit the graph directly (not just embeddings)

tax_embedding_technical_overview.md - Technical Overview: Graph-Based Contrastive Learning for Tax Document Retrieval

What We Learned

1. Tax documents are graphs, not text

IRS forms and instructions are inherently **relational**:

- Meaning lives in **edges** (“includes”, “exception to”, “see also”), not just nodes.
- Any system that treats them as independent chunks will *necessarily* miss meaning.

👉 **Insight:** Retrieval quality is capped unless document relationships are modeled explicitly.

2. Naive RAG fails because it destroys semantic dependencies

Chunking by box, section, or token count:

- Breaks parent-child relationships
- Loses cross-box logic (e.g., *Box 1a includes 1b and 2e*)
- Cannot answer comparative or concept-level questions

👉 **Insight:** The failure is structural, not just model quality.

3. Embedding models only learn what training pairs teach

Out-of-the-box embeddings:

- Don't know that "Box 1a" and "Box 1b" are related
- Treat similar-looking but unrelated boxes as close
- Miss negation and exception semantics

👉 **Insight:** Retrieval improves only if we **explicitly teach relationships** during training.

4. Graphs are the right abstraction for both understanding and training

By modeling chunks as nodes and references as edges, we unlock:

- High-quality **positive pairs** (graph neighbors)
- **Hard negatives** (same doc, similar surface form, no edge)
- Hierarchical learning (section → subsection)
- Cross-document alignment (filer ↔ recipient ↔ form)

👉 **Insight:** The graph becomes a *supervisor* for contrastive learning.

5. Contrastive learning works best when pairs encode meaning, not proximity

Effective pairs come from:

- `includes`, `exception_to`, `same_field`, `defines`
- Cross-document equivalence
- Concept ↔ atomic grounding

Ineffective pairs come from:

- Page proximity
- Random same-section sampling

👉 **Insight:** Pair *quality* matters more than pair *volume*.

6. Conceptual chunks are necessary—but must be constrained

Atomic chunks are faithful but fragmented.

Conceptual chunks:

- Enable concept-level queries (“What are qualified dividends?”)
- Improve recall and ranking dramatically

But only if:

- They are **source-grounded**
- Fully traceable (`derived_from`)
- Strictly validated for faithfulness

👉 **Insight:** Conceptual synthesis is powerful **only when auditable**.

7. Graph distance is a better negative signal than randomness

Hard negatives are most effective when they are:

- Same form
- Same structural level
- Numerically or lexically similar
- But **graph-distant**

Example:

- Box 1b vs Box 2a (both dividend amounts, different meaning)

👉 **Insight:** Graph topology gives us principled hard negatives for free.

8. Training must balance multiple relationship types

No single pair type is sufficient:

- Query—passage teaches retrieval
- Graph similarity teaches semantic structure
- Cross-doc pairs teach alignment
- Hierarchical pairs teach abstraction

- Concept↔atomic pairs teach grounding

👉 **Insight:** Embedding quality is a function of **relationship diversity**, not just data size.

9. Retrieval should be concept-aware, not chunk-only

Best-performing strategies:

- Retrieve **conceptual chunks first**, then expand to atomics
- Or retrieve both and re-rank jointly

Worst-performing:

- Atomic-only retrieval with large **k**

👉 **Insight:** Concepts should guide retrieval, not just embeddings.

10. Success is measurable—and looks different

Improvement isn't just higher similarity scores:

- Concept queries retrieve **complete semantic coverage**
- Comparative queries retrieve **both sides + relationship**
- Edge cases retrieve **rules + exceptions**

👉 **Insight:** Metrics like **concept coverage** matter as much as MRR.

The Core Takeaway (One Sentence)

To make tax documents retrievable, we must train embedding models on the same relational structure that humans use to understand them—and graphs are the missing supervision layer.

If you want, next we can:

- Pressure-test this against another IRS form (e.g., 1099-INT or W-2)
- Translate this into a reusable “Graph-RAG” pattern

- Or define what *not* to model in the graph to keep it scalable

Here's a **clean synthesis of what we learned** from this Living Strategy — not a restatement, but the distilled insights that now guide decisions.

1. The Core Insight: Tax Retrieval Is a *Structure* Problem First, Not an Embedding Problem

- Tax documents are **not flat text**. They are:
 - Hierarchical (form → section → box → paragraph)
 - Cross-referential (“Box 1a includes 1b”)
 - Precision-sensitive (“ $2e \neq 2f$ ”)
- Any system that ignores structure will plateau quickly, regardless of model quality.

Learning:

Retrieval quality is bottlenecked by **structural fidelity**, not model capacity.

This justifies:

- Knowledge graphs
 - Registries
 - Deterministic extraction
 - Provenance-first design
-

2. Fine-Tuning Is a *Conditional Bet*, Not a Given

You explicitly framed fine-tuning as **falsifiable**, which is a major learning.

- Strong baselines today (contextual prefixes + BM25 + rerank) may already solve 70–80%
- Fine-tuning is only justified if it **materially beats** that baseline (+10–15% Recall@5)

Learning:

We are not “building embeddings”; we are running an experiment with a clear kill switch.

This prevents sunk-cost bias and keeps the architecture honest.

3. The Knowledge Graph’s Real Value Is *Training Signal Control*, Not Just Retrieval

A subtle but critical realization:

- The KG is not primarily for GraphRAG-style summarization
- Its highest leverage use is:
 - Generating **trusted contrastive pairs**
 - Filtering false negatives
 - Weighting training loss by edge confidence

Learning:

The KG is a *data quality machine* more than a retrieval feature.

Retrieval benefits are incremental; training benefits are foundational.

4. Granularity Must Be Multi-Level — There Is No Single “Right Chunk Size”

You resolved a classic tension cleanly:

- Training wants **bigger, semantically rich units**
- Retrieval wants **small, precise units**

Learning:

One graph, multiple granularities, different query patterns.

Hierarchy is the unifying abstraction:

- Anchors for training
- Paragraphs for retrieval
- Sentences only when density demands it

This avoids building parallel systems.

5. Confidence + Provenance Is the Only Scalable Way to Use LLMs Safely

The document makes this explicit and operational:

- LLMs are **proposers, not authorities**
- Every decision must:
 - Cite evidence
 - Emit confidence
 - Be rejectable

Learning:

Debuggability beats raw capability.

This is why:

- Edge confidence tiers matter
- LLM-as-judge is gated
- “Unsupported claims” fail hard

This turns LLM usage from magic into engineering.

6. Most “Hard” Problems Are Actually *Routing* Problems

A key emerging realization in the Open Questions:

- Many queries don’t need retrieval at all:
 - “Explain Box 2a”
 - “What goes in Box 1b”
- These are deterministic alias → canonical ID lookups

Learning:

If we don’t build a router, we’ll overuse retrieval and underperform.

This reframes part of the system from “search” to “structured navigation.”

7. Precision Failures Are Inevitable With Single-Vector Models — So We Hedge

You didn’t commit blindly to embeddings:

- You acknowledged known failure modes (2e vs 2f)
- You built in:
 - BM25
 - Cross-encoder reranking
 - ColBERT / SPLADE benchmarks

Learning:

Architectural redundancy is intentional, not indecision.

This keeps options open and prevents lock-in.

8. Evaluation Is a First-Class System, Not an Afterthought

Another strong takeaway:

- Evaluation is:
 - Stratified
 - Precision-aware
 - Multi-hop aware
- Near-miss negatives prevent false confidence

Learning:

If the eval can’t fail the system, the system isn’t real.

This ensures progress is measurable and meaningful.

9. Phasing Is About *Risk Burn-Down*, Not Features

Phase A isn't about shipping features — it's about validating assumptions:

- Can we extract reliably?
- Can we build a clean graph?
- What is the real baseline ceiling?

Learning:

Each phase answers a “should we continue?” question.

That's a mature systems mindset.

10. The Meta-Learning: This Is an Anti-Overengineering Strategy

Despite the complexity, the document is deeply **anti-overbuild**:

- Multiple explicit pivot points
- Cheap baselines compared
- Alternatives (CoBERT, routing, cheap extraction) kept alive

Learning:

Complexity is earned only if numbers demand it.

One-Line Summary

We learned that winning tax document intelligence is about controlling structure, trust, and evaluation—not about betting on smarter models—and that every expensive component must justify itself against a strong, simpler baseline.

If you want, next we can:

- Turn this into a **decision tree** (“If metric X < Y, do Z”)
- Extract a **1-page exec version**

- Or map this to a reference architecture diagram

`Tax_rag_living_strategy.md`:

Here's a clean synthesis of what we learned from this Living Strategy — not a restatement, but the distilled insights that now guide decisions.

1. The Core Insight: Tax Retrieval Is a *Structure* Problem First, Not an Embedding Problem

- Tax documents are not flat text. They are:
 - Hierarchical (form → section → box → paragraph)
 - Cross-referential (“Box 1a includes 1b”)
 - Precision-sensitive (“ $2e \neq 2f$ ”)
- Any system that ignores structure will plateau quickly, regardless of model quality.

Learning:

Retrieval quality is bottlenecked by structural fidelity, not model capacity.

This justifies:

- Knowledge graphs
- Registries
- Deterministic extraction
- Provenance-first design

2. Fine-Tuning Is a *Conditional Bet*, Not a Given

You explicitly framed fine-tuning as falsifiable, which is a major learning.

- Strong baselines today (contextual prefixes + BM25 + rerank) may already solve 70–80%
- Fine-tuning is only justified if it materially beats that baseline (+10–15% Recall@5)

Learning:

We are not “building embeddings”; we are running an experiment with a clear kill switch.

This prevents sunk-cost bias and keeps the architecture honest.

3. The Knowledge Graph’s Real Value Is *Training Signal Control*, Not Just Retrieval

A subtle but critical realization:

- The KG is not primarily for GraphRAG-style summarization
- Its highest leverage use is:
 - Generating trusted contrastive pairs
 - Filtering false negatives
 - Weighting training loss by edge confidence

Learning:

The KG is a *data quality machine* more than a retrieval feature.

Retrieval benefits are incremental; training benefits are foundational.

4. Granularity Must Be Multi-Level — There Is No Single “Right Chunk Size”

You resolved a classic tension cleanly:

- Training wants bigger, semantically rich units
- Retrieval wants small, precise units

Learning:

One graph, multiple granularities, different query patterns.

Hierarchy is the unifying abstraction:

- Anchors for training

- Paragraphs for retrieval
- Sentences only when density demands it

This avoids building parallel systems.

5. Confidence + Provenance Is the Only Scalable Way to Use LLMs Safely

The document makes this explicit and operational:

- LLMs are proposers, not authorities
- Every decision must:
 - Cite evidence
 - Emit confidence
 - Be rejectable

Learning:

Debuggability beats raw capability.

This is why:

- Edge confidence tiers matter
- LLM-as-judge is gated
- “Unsupported claims” fail hard

This turns LLM usage from magic into engineering.

6. Most “Hard” Problems Are Actually *Routing* Problems

A key emerging realization in the Open Questions:

- Many queries don’t need retrieval at all:
 - “Explain Box 2a”
 - “What goes in Box 1b”
- These are deterministic alias → canonical ID lookups

Learning:

If we don't build a router, we'll overuse retrieval and underperform.

This reframes part of the system from "search" to "structured navigation."

7. Precision Failures Are Inevitable With Single-Vector Models — So We Hedge

You didn't commit blindly to embeddings:

- You acknowledged known failure modes (2e vs 2f)
- You built in:
 - BM25
 - Cross-encoder reranking
 - ColBERT / SPLADE benchmarks

Learning:

Architectural redundancy is intentional, not indecision.

This keeps options open and prevents lock-in.

8. Evaluation Is a First-Class System, Not an Afterthought

Another strong takeaway:

- Evaluation is:
 - Stratified
 - Precision-aware
 - Multi-hop aware
- Near-miss negatives prevent false confidence

Learning:

If the eval can't fail the system, the system isn't real.

This ensures progress is measurable and meaningful.

9. Phasing Is About *Risk Burn-Down*, Not Features

Phase A isn't about shipping features — it's about validating assumptions:

- Can we extract reliably?
- Can we build a clean graph?
- What is the real baseline ceiling?

Learning:

Each phase answers a “should we continue?” question.

That's a mature systems mindset.

10. The Meta-Learning: This Is an Anti-Overengineering Strategy

Despite the complexity, the document is deeply anti-overbuild:

- Multiple explicit pivot points
- Cheap baselines compared
- Alternatives (CoBERT, routing, cheap extraction) kept alive

Learning:

Complexity is earned only if numbers demand it.

One-Line Summary

We learned that winning tax document intelligence is about controlling structure, trust, and evaluation—not about betting on smarter models—and that every expensive component must justify itself against a strong, simpler baseline

